

Notebook

June 29, 2025

```
[1]: import os
import shutil
import random
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

# Reproduzierbarkeit sicherstellen
tf.random.set_seed(42)
np.random.seed(42)
random.seed(42)
```

```
2025-06-29 22:22:03.807622: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-06-29 22:22:03.813648: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-06-29 22:22:03.829860: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1751228523.862913    20679 cuda_dnn.cc:8579] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1751228523.871574    20679 cuda_blas.cc:1407] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
W0000 00:00:1751228523.894053    20679 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
target more than once.
W0000 00:00:1751228523.894085    20679 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
```

target more than once.

W0000 00:00:1751228523.894087 20679 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.

W0000 00:00:1751228523.894089 20679 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.

2025-06-29 22:22:03.903361: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
[2]: DATA_DIR = "../data/raw"
IMG_SIZE = (200, 300)
BATCH_SIZE = 32

classes = ['paper', 'rock', 'scissors']
class_counts = {}

print("=== DATENÜBERSICHT ===")
for cls in classes:
    cls_path = os.path.join(DATA_DIR, cls)
    images = [f for f in os.listdir(cls_path) if f.endswith('.png')]
    class_counts[cls] = len(images)
    print(f"{cls}: {len(images)} Bilder")

total_images = sum(class_counts.values())
print(f"\nGesamtanzahl: {total_images} Bilder")
print(f"Durchschnitt pro Klasse: {total_images/3:.1f}")

plt.figure(figsize=(8,5))
bars = plt.bar(classes, [class_counts[cls] for cls in classes],
               color=['lightblue', 'lightcoral', 'lightgreen'])
plt.title("Klassenverteilung im Dataset")
plt.ylabel("Anzahl Bilder")
for bar, count in zip(bars, [class_counts[cls] for cls in classes]):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 5,
             str(count), ha='center', va='bottom')
plt.show()
```

=== DATENÜBERSICHT ===


paper: 712 Bilder

rock: 726 Bilder

scissors: 750 Bilder

Gesamtanzahl: 2188 Bilder

Durchschnitt pro Klasse: 729.3



```
[3]: DATA_DIR = "../data/raw"
IMG_SIZE = (200, 300)
BATCH_SIZE = 32
SEED = 42
VAL_SPLIT = 0.30

train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATA_DIR,
    validation_split=VAL_SPLIT,
    subset="training",
    seed=SEED,
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
val_test_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATA_DIR,
    validation_split=VAL_SPLIT,
    subset="validation",
    seed=SEED,
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
```

```

val_batches = tf.data.experimental.cardinality(val_test_ds)
val_ds = val_test_ds.take(val_batches // 2)
test_ds = val_test_ds.skip(val_batches // 2)

print(f"Train-Batches: {tf.data.experimental.cardinality(train_ds)}")
print(f"  Val-Batches: {tf.data.experimental.cardinality(val_ds)}")
print(f" Test-Batches: {tf.data.experimental.cardinality(test_ds)}")

```

Found 2188 files belonging to 3 classes.

Using 1532 files for training.

E0000 00:00:1751228530.318254 20679 cuda_executor.cc:1228] INTERNAL: CUDA Runtime error: Failed call to cudaGetRuntimeVersion: Error loading CUDA libraries. GPU will not be used.: Error loading CUDA libraries. GPU will not be used.

W0000 00:00:1751228530.319164 20679 gpu_device.cc:2341] Cannot dlopen some GPU libraries. Please make sure the missing libraries mentioned above are installed properly if you would like to use GPU. Follow the guide at <https://www.tensorflow.org/install/gpu> for how to download and setup the required libraries for your platform.

Skipping registering GPU devices...

Found 2188 files belonging to 3 classes.

Using 656 files for validation.

Train-Batches: 48

Val-Batches: 10

Test-Batches: 11

```


[4]: data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.2),
    layers.RandomZoom(0.1),
])

plt.figure(figsize=(8,8))
for images, _ in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3,3,i+1)
        aug_img = data_augmentation(tf.expand_dims(images[i], 0))
        plt.imshow(aug_img[0].numpy().astype("uint8"))
        plt.axis("off")
plt.suptitle("Live-Augmentation (300×200)")
plt.tight_layout()
plt.show()

```

2025-06-29 22:22:12.277385: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

Live-Augmentation (300×200)



```
[5]: def build_cnn(input_shape=(200,300,3), n_classes=3):  
    model = models.Sequential([  
        layers.Input(shape=input_shape),  
  
        layers.Conv2D(32, (3,3), activation='relu', padding='same'),  
        layers.MaxPooling2D((2,2)),  
  
        layers.Conv2D(64, (3,3), activation='relu', padding='same'),  
        layers.MaxPooling2D((2,2)),  
  
        layers.Conv2D(128, (3,3), activation='relu', padding='same'),  
        layers.MaxPooling2D((2,2)),  
  
        layers.GlobalAveragePooling2D(),  
        layers.Dropout(0.5),
```

```

        layers.Dense(64, activation='relu'),
        layers.Dense(n_classes, activation='softmax'),
    ])
    return model

model = build_cnn()
model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 200, 300, 32)	896
max_pooling2d (MaxPooling2D)	(None, 100, 150, 32)	0
conv2d_1 (Conv2D)	(None, 100, 150, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 50, 75, 64)	0
conv2d_2 (Conv2D)	(None, 50, 75, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 25, 37, 128)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 128)	0
dropout (Dropout)	(None, 128)	0
dense (Dense)	(None, 64)	8,256
dense_1 (Dense)	(None, 3)	195

Total params: 101,699 (397.26 KB)

Trainable params: 101,699 (397.26 KB)

Non-trainable params: 0 (0.00 B)

```

[6]: model.compile(
      optimizer='adam',

```

```

        loss='sparse_categorical_crossentropy',
        metrics=['accuracy']
    )

    callbacks = [
        EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True),
        ModelCheckpoint("../models/rps_best.h5", save_best_only=True)
    ]

```

```

[7]: history = model.fit(
        train_ds,
        validation_data=val_ds,
        epochs=15,
        callbacks=callbacks
    )

```

Epoch 1/15

48/48 0s 932ms/step -
accuracy: 0.3647 - loss: 8.0571

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

48/48 49s 987ms/step -
accuracy: 0.3649 - loss: 7.9685 - val_accuracy: 0.4781 - val_loss: 1.0648

Epoch 2/15

48/48 0s 935ms/step -
accuracy: 0.4106 - loss: 1.0788

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

48/48 47s 988ms/step -
accuracy: 0.4110 - loss: 1.0780 - val_accuracy: 0.5688 - val_loss: 0.8289

Epoch 3/15

48/48 0s 943ms/step -
accuracy: 0.6371 - loss: 0.8010

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

48/48 48s 994ms/step -

accuracy: 0.6373 - loss: 0.8008 - val_accuracy: 0.9094 - val_loss: 0.5201
Epoch 4/15

48/48 0s 936ms/step -

accuracy: 0.7965 - loss: 0.5493

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

48/48 47s 987ms/step -

accuracy: 0.7964 - loss: 0.5489 - val_accuracy: 0.9219 - val_loss: 0.3425

Epoch 5/15

48/48 0s 946ms/step -

accuracy: 0.8644 - loss: 0.3710

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

48/48 48s 997ms/step -

accuracy: 0.8646 - loss: 0.3707 - val_accuracy: 0.9594 - val_loss: 0.1904

Epoch 6/15

48/48 48s 1s/step -

accuracy: 0.8898 - loss: 0.3013 - val_accuracy: 0.9563 - val_loss: 0.1908

Epoch 7/15

48/48 0s 938ms/step -

accuracy: 0.9066 - loss: 0.2652

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

48/48 47s 989ms/step -

accuracy: 0.9068 - loss: 0.2650 - val_accuracy: 0.9500 - val_loss: 0.1893

Epoch 8/15

48/48 47s 991ms/step -

accuracy: 0.9036 - loss: 0.2780 - val_accuracy: 0.8375 - val_loss: 0.3559

Epoch 9/15

48/48 0s 943ms/step -

accuracy: 0.8812 - loss: 0.2988

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,

'my_model.keras')`.

48/48 48s 996ms/step -
accuracy: 0.8818 - loss: 0.2980 - val_accuracy: 0.9719 - val_loss: 0.1388
Epoch 10/15

48/48 0s 954ms/step -
accuracy: 0.9403 - loss: 0.1961

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

48/48 48s 1s/step -
accuracy: 0.9405 - loss: 0.1960 - val_accuracy: 0.9719 - val_loss: 0.0969
Epoch 11/15

48/48 47s 977ms/step -
accuracy: 0.9499 - loss: 0.1865 - val_accuracy: 0.9594 - val_loss: 0.1389
Epoch 12/15

48/48 0s 945ms/step -
accuracy: 0.9480 - loss: 0.1682

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

48/48 48s 995ms/step -
accuracy: 0.9480 - loss: 0.1683 - val_accuracy: 0.9781 - val_loss: 0.0905
Epoch 13/15

48/48 47s 973ms/step -
accuracy: 0.9656 - loss: 0.1611 - val_accuracy: 0.9563 - val_loss: 0.1233
Epoch 14/15

48/48 47s 973ms/step -
accuracy: 0.9482 - loss: 0.1550 - val_accuracy: 0.9656 - val_loss: 0.1192
Epoch 15/15

48/48 0s 915ms/step -
accuracy: 0.9403 - loss: 0.1661

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.


48/48 46s 964ms/step -
accuracy: 0.9403 - loss: 0.1665 - val_accuracy: 0.9844 - val_loss: 0.0741

```
[8]: hist = history.history
plt.figure(figsize=(10,4))

plt.subplot(1,2,1)
plt.plot(hist['loss'], label='Train Loss')
plt.plot(hist['val_loss'], label='Val Loss')
plt.legend(); plt.title("Loss")

plt.subplot(1,2,2)
plt.plot(hist['accuracy'], label='Train Acc')
plt.plot(hist['val_accuracy'], label='Val Acc')
plt.legend(); plt.title("Accuracy")

plt.tight_layout()
plt.show()
```



```
[11]: y_pred_probs_list = []
y_true_list = []

for images, labels in test_ds:
    preds = model.predict_on_batch(images)
    y_pred_probs_list.append(preds)
    y_true_list.append(labels.numpy())

y_pred_probs = np.concatenate(y_pred_probs_list, axis=0)
y_true = np.concatenate(y_true_list, axis=0)

y_pred = np.argmax(y_pred_probs, axis=1)

class_names = train_ds.class_names
```

```

print("=== Classification Report ===")
print(classification_report(y_true, y_pred, target_names=class_names))

cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix')
plt.ylabel('Wahre Klasse (True)')
plt.xlabel('Vorhergesagte Klasse (Predicted)')
plt.show()

```

```

=== Classification Report ===

```

	precision	recall	f1-score	support
paper	0.97	0.96	0.97	101
rock	1.00	0.99	1.00	115
scissors	0.97	0.98	0.98	120
accuracy			0.98	336
macro avg	0.98	0.98	0.98	336
weighted avg	0.98	0.98	0.98	336

```

2025-06-29 23:25:48.394135: I tensorflow/core/framework/local_rendezvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

```

